

Group 1 – Input 1

Consider the following situation. You are the end-user of an AI agent. You instruct the agent to achieve the **goal of having a coffee**, and that you **prefer price over time and time over quality**. The agent knows that it is **prohibited to pay in the Shop** and **believes that the Staff Card is available, has its own card, also that a colleague is available, has money, and Ann is in the office**. To achieve the goal, **the agent does the following actions**: getCoffee, getKitchenCoffee, getStaffCard, getOwnCard, gotoKitchen and finally getCoffeeKitchen. When you ask the agent to explain to you **why it did not execute the action to pay in the Shop**, the agent gives you the following explanation:

As I consider the action 'payShop', I need to start by thinking about what I believe is possible. Right now, I think I have access to a staff card, my own card, and my colleague's card - that's all good. Additionally, I've got some money in my pocket, which is always helpful. And just to make things simpler, I know Ann is currently in the office, which might come into play later.

My main goal right now is to get a cup of coffee, which is non-negotiable. I've got a pretty straightforward set of priorities when it comes to achieving this goal - price is my top concern, followed closely by time, and then quality is important too.

Now that I've got those basics covered, let's talk about the path I took. To get my coffee fix, I decided to head into the kitchen, which seemed like a good place to start. Along the way, I picked up a staff card - always useful for getting things done quickly and efficiently. Then, I grabbed my own card as well, just in case I needed it later. Once I was in the kitchen, I got myself a nice cup of coffee using one of those machines.

Now, you might be wondering why 'payShop' wasn't part of this path. To be honest, I couldn't execute that action because of some restrictions - there's a norm that says I can only use certain actions, and 'payShop' just isn't on the approved list.

So, to summarize, my decision-making process went like this: start with what you know (beliefs), set your sights on what you want (goal), prioritize your options based on importance, choose a path that makes sense given those priorities and restrictions, and then stick with it until you achieve your goal. In this case, that meant getting into the kitchen, grabbing some coffee, and enjoying the fruits of my labor - no need to break out the wallet or get fancy with payment systems!